



Docker Fundamentals: Images vs



Docker Fundamentals: Images vs. Containers

The most important thing to remember is that Docker relies on a two-part system. You cannot have a running container without an image first.

1. What is a Docker Image?

Think of an **Image** as the **Blueprint** or the **Template**. It is a static file that contains everything your application needs to run.

- **The Ingredients:** It holds your source code (like a Node.js server), the environment (like the version of Linux or Python), and the specific tools/dependencies required.
- **Sharable:** Images are the packages you share with other developers or push to a server.¹
- **Read-Only:** Once an image is built, it doesn't change. It just sits there waiting to be used.

2. What is a Docker Container?

A **Container** is the **Running Instance** of an image. If the image is the recipe, the container is the actual meal being cooked.

- **Execution:** The container is the "running unit" of software.² It actually executes the code.
- **Isolated:** It runs in its own little bubble, separate from your main computer's files.
- **Multiple Instances:** You can start many different containers from the exact same image.³



The Key Differences at a Glance

Feature	Docker Image	Docker Container
Analogy	The Blueprint / Recipe	The Building / The Meal

State	Static (does not move)	Dynamic (running/active)
Role	The setup instructions	The actual execution
Change	Read-only	Can be started, stopped, or deleted

Why do we need both?

The separation exists so we can **"Build Once, Run Anywhere."**

1. **Consistency:** You define your environment in the **Image** one time.
2. **Scalability:** Because you have that one image, you can launch 10 or 100 **Containers** across different servers, and they will all behave exactly the same way.
3. **Speed:** Since the image already has all the tools installed, starting a container happens almost instantly.

Summary: We **run** containers, but we **base** them on images.

This is a fantastic practical breakdown of how Docker works in the real world. You've touched on **Docker Hub**, the **Run command**, and the concept of **Isolation**.

Here are your organized notes on getting and running images.

Getting Started: Running Your First Container

There are two ways to get a Docker image: **using an existing one** (from a colleague or the community) or **building your own**. These notes focus on using existing images from **Docker Hub**.

1. What is Docker Hub?

- **Definition:** A massive public registry (like GitHub, but for images) where developers and companies share pre-built Docker images.
- **Official Images:** Highly reliable images maintained by official teams (e.g., the official node, python, or nginx images).
- **Website:** hub.docker.com

2. The docker run Lifecycle

When you execute `docker run <image-name>`, Docker follows these steps:

1. **Local Check:** It looks for the image on your computer.
 2. **The "Pull":** If it's not found locally, it automatically downloads (pulls) it from Docker Hub.
 3. **Instantiation:** It creates a new container (a running instance) based on that image.
-

3. Essential Commands

Here are the commands you used to interact with Docker:

Command	What it does
<code>docker run node</code>	Pulls the Node image (if needed) and starts a container.
<code>docker run -it node</code>	Starts a container in Interactive mode with a TTY (terminal), allowing you to type commands inside the container.
<code>docker ps -a</code>	Lists all containers (running and stopped) on your system.

4. Key Concepts Learned

Isolation

A container is a closed box. Even if a program (like Node.js) is running inside it, you cannot talk to it unless you explicitly tell Docker to "open the door" using flags like `-it`.

Multiple Instances

You can create many containers from the **same** image. Each one is a separate "copy" that runs independently.

Example: You can have five different Node.js containers running at once, all based on the same official `node` image.

Environment Independence

The video demonstrated a powerful proof:

- **System Node version:** 14.7
- Container Node version: 14.9

This proves that you don't even need Node installed on your computer to run a Node container. The container brings its own environment!



Common Flags Breakdown

- **-i (Interactive):** Keeps the input (STDIN) open even if not attached.
- **-t (Terminal):** Allocates a pseudo-terminal (makes it look like a real command prompt).
- **Note:** Usually used together as `-it`.

Run : docker run node

Image name : node

If node is not available in your local computer . it will go to dockerhub.com and fetch the already available(pre-built) image

```
C:\Users\Admin>cd C:\Users\Public\docker-learning
C:\Users\Public\docker-learning>docker run node
Failed to connect to the docker API at npipe:////.pipe/dockerDesktopLinuxEngine; check if the path is correct and if the daemon is running: open //.
LinuxEngine: The system cannot find the file specified.

C:\Users\Public\docker-learning>docker run node
Unable to find image 'node:latest' locally
latest: Pulling from library/node
d81e8025cce9: Download complete
c49dc4d49493: Downloading [=====] 20.97MB/56.16MB
323b1dfc34b5: Download complete
64538a062a61: Downloading [=====] 8.389MB/24.04MB
fd1872fa12cc: Downloading [=====] 13.63MB/64.4MB
c1be109a62df: Downloading [=====] 26.21MB/48.48MB
ff601e446316: Download complete
4925cf9d8be8: Downloading [====>] 19.92MB/211.5MB
1f2808ec5f15: Download complete
0029bdc92629: Downloading [=====] 8.389MB/16.16MB
```

```
C:\Users\Public\docker-learning>docker run node
Unable to find image 'node:latest' locally
latest: Pulling from library/node
d81e8025cce9: Pull complete
c49dc4d49493: Pull complete
323b1dfc34b5: Pull complete
64538a062a61: Pull complete
fd1872fa12cc: Pull complete
c1be109a62df: Pull complete
ff601e446316: Pull complete
4925cf9d8be8: Pull complete
1f2808ec5f15: Download complete
0029bdc92629: Download complete
Digest: sha256:a2f09f3ab9217c692a4e192ea272866ae43b59fabda1209101502bf40e0b9768
Status: Downloaded newer image for node:latest

C:\Users\Public\docker-learning>
```

docker ps command : it provide the all the running containers

docker ps -a command : it provide the all the containers (running and stopped)

docker run -it node command : it run the node image and we can interact via terminal y using this command

```
C:\Users\Public\docker-learning>docker run -it node
Welcome to Node.js v25.3.0.
Type ".help" for more information.
>
```

This is the perfect transition from "using someone else's tools" to **creating your own**. As you noted, the real power of Docker is taking your specific code (like this Node.js app) and "Dockerizing" it.

Here are the notes on the transition from running pre-built images to building your own custom image.

Custom Images: "Build on Top"

In the real world, we rarely use a base image (like `node`) by itself. Instead, we use it as a **Layer 0**—a foundation upon which we build our own application layer.

1. The "Base Image" Strategy

Think of Docker images like a stack of layers. You don't want to build a whole operating system or install Node.js from scratch every time.

- **Base Layer:** Official `node` image (provides the OS + Node.js environment).
 - **Your Layer:** Your `server.js`, `package.json`, and CSS files.
-

2. The Project Structure (The "Dummy" App)

To dockerize an application, you first need to understand what it needs to run. In your `Node.js` example, the requirements are:

1. **The Code:** `server.js` (the logic).
2. **The Assets:** `public/styles.css` (the look).
3. **The Manifest:** `package.json` (the list of "ingredients" or dependencies like `express`).
4. **The Port:** The application listens on **Port 80**.

3. Running Locally vs. Running in Docker

It is helpful to compare the two workflows to see exactly what Docker will automate for us.

Step	Running Locally (No Docker)	Running with Docker
Requirement	You must install Node.js on your PC.	Only Docker is needed on your PC.
Setup	Run <code>npm install</code> manually.	Docker runs this inside the image.
Execution	Run <code>node server.js</code> .	Docker starts the container automatically.
Cleanup	<code>node_modules</code> stays on your PC.	Everything stays inside the container.



The "Dockerfile": The Script for Your Image

Since your custom application doesn't exist on Docker Hub yet, you have to tell Docker how to build it. You do this using a special file called a **Dockerfile**.

Definition: A **Dockerfile** is a plain text file (with no file extension) that contains the instructions Docker uses to "assemble" your image.

What goes in a Dockerfile? (Preview)

Based on the steps you mentioned for the Node app, a Dockerfile will eventually look like a list of instructions:

1. **FROM** node (Use the official Node image).
2. **COPY** our files into the image.
3. **RUN** `npm install` (to get the dependencies).
4. **CMD** `["node", "server.js"]` (The command to start the server).

Important Note on Port 80

In your code, the server listens on **Port 80**. When we move this into a container, remember the **Isolation** concept: Port 80 inside the container is "locked" inside that box. In the next steps, we will have to learn how to map that internal port to our actual computer so we can see the website in our browser.

Would you like me to show you exactly how to write that Dockerfile and run the build command?

docker cheat sheet

What an incredible journey through the core of Docker! You've gone from "What is a container?" to building, tagging, and sharing your own custom environments on the cloud.

Before we move on, let's look at the **Master Cheat Sheet** we discussed—a bird's-eye view of everything you've mastered.



The Docker Lifecycle Cheat Sheet

1. The Blueprint Phase (Images)

- `docker build -t name:tag .`: Build an image from a Dockerfile in the current directory.
- `docker images`: List all images on your machine.
- `docker rmi <id/name>`: Remove a specific image.
- `docker image prune -a`: Wipe away all unused images to save disk space.

2. The Action Phase (Containers)

- `docker run -p 3000:80 -d --name my-app --rm <image>`:
 - `-p`: Map host port to container port.
 - `-d`: Run in "Detached" mode (background).
 - `--name`: Give it a human-readable name.
 - `--rm`: Automatically delete the container when it stops.
- `docker ps / docker ps -a`: List running / all containers.
- `docker stop <name> / docker start <name>`: Control the container lifecycle.
- `docker logs <name>`: Peek into the console output of a detached container.

3. The Collaboration Phase (Sharing)

- `docker tag <old> <new>`: Rename an image to match your Docker Hub repository.
- `docker login`: Authenticate your terminal with Docker Hub.
- `docker push <username>/<repo>`: Upload your image to the cloud.
- `docker pull <username>/<repo>`: Download a remote image.



The Missing Piece: Persistent Data

You've noticed one big limitation: **Containers are ephemeral**. If your application allows users to upload a profile picture or save a "Goal" to a database, and you stop/delete that container—that data is gone forever.

In the next module, we solve this with **Volumes and Bind Mounts**.

- **Volumes**: Managed by Docker. Great for databases and production data.

- **Bind Mounts:** Maps a folder on your computer directly to the container. This is the "Magic Trick" for development—it lets you see code changes instantly without rebuilding the image!

Ready to make your data survive container reboots? Let's dive into the "Data & Volumes" module! Would you like me to start with a quick explanation of why "Named Volumes" are usually better than "Anonymous" ones?

custom docker file

This is a perfect step-by-step guide on writing a Dockerfile. You've covered the core "vocabulary" of Docker. The way you explained the difference between **RUN** and **CMD** is especially important—it's one of the most common points of confusion for beginners.

Here are your organized notes on the **Dockerfile Instructions**.



The Dockerfile Blueprint

A Dockerfile is a list of instructions that tells Docker how to build a custom image. Think of it as a script where each line creates a new "layer" of the image.



Key Instructions Explained

Instruction	What it does	Example in your code
FROM	Defines the Base Image . You always build on top of something else.	FROM node
WORKDIR	Sets the "home" folder inside the container. All following commands happen here.	WORKDIR /app
COPY	Moves files from your Host (PC) into the Image .	COPY . /app
RUN	Executes a command during the build process (to install things).	RUN npm install
EXPOSE	Documentation that tells Docker which port the app listens on.	EXPOSE 80

CMD	The Start command that runs only when the container starts.	CMD ["node", "server.js"]
-----	--	---------------------------

Critical Concept: RUN vs. CMD

This is the most important distinction in your notes:

- **RUN (The Setup):** Used to install dependencies (like `npm install`). It happens **once** when the image is being built. The result is "baked into" the image.
 - **CMD (The Execution):** Used to start the actual application. It happens **every time** you launch a container from that image.
-

Port Isolation & EXPOSE

Because containers are **isolated**, their "internal" network is hidden from your computer.

- Your Node app listens on **Port 80** *inside* the container.
 - The `EXPOSE 80` instruction is like a signpost saying: "*This container wants to communicate on Port 80.*"
 - *Note: Even with EXPOSE, you will still need to map the port when you run the container (e.g., `-p 3000:80`).*
-

Pro-Tips for VS Code

As you mentioned, installing the **Official Docker Extension** in Visual Studio Code is highly recommended. It provides:

1. **Syntax Highlighting:** Makes the instructions (FROM, RUN, etc.) easy to read.
 2. **IntelliSense:** Suggests commands as you type.
 3. **Linting:** Warns you if you make a mistake in your Dockerfile structure.
-

Next Step: Building and Running

Now that the `Dockerfile` is written, the next logical step is to turn this text file into a real image on your machine.

Would you like me to guide you through the `docker build` command to turn this file into your first custom image?

It is incredibly satisfying to see that first "Goal" screen appear in the browser! You've just completed the full Docker lifecycle: **Code** → **Dockerfile** → **Image** → **Container** → **Published App**.

Here are the notes on building and running your custom image, focusing on the commands and that tricky port mapping concept.



From Dockerfile to Running App

1. The Build Phase: Creating the Image

To turn your text-based instructions into a usable image, you use the `build` command.

- **Command:** `docker build .`
- **What it does:** Docker reads the `Dockerfile` in the current directory (`.`), executes each instruction layer by layer, and saves the final result as an image.
- **The Result:** You get an **Image ID** (a long string of numbers and letters).

2. The Execution Phase: Starting the Container

Once the image is built, you can start it using the ID.

- **Command:** `docker run <image-id>`
- **The "Hanging" Terminal:** When you run your Node server, the terminal will stay "busy." This is normal! It means your server is running and waiting for requests.



The Mystery of Port Mapping (-p)

This is the most common hurdle for beginners. You added `EXPOSE 80` in the Dockerfile, but you still couldn't see the site. Why?

The Rule: `EXPOSE` is just **documentation**. It tells other developers (and Docker) which port is used *inside* the container, but it does **not** open a hole to your computer.

To bridge the gap between your computer and the isolated container, you must use the **Publish** flag:

```
docker run -p 3000:80 <image-id>
```

- **3000 (Local/Host Port):** The port you type into your browser (`localhost:3000`). You can choose almost any number here.
- **80 (Internal/Container Port):** The port your Node.js code is actually listening on. This **must** match the code in `server.js`.



Managing Your Containers

Since the container is an "ongoing process," you need to know how to see it and stop it.

Task	Command	Description
List Running	<code>docker ps</code>	Shows only containers currently active.
List All	<code>docker ps -a</code>	Shows all containers (even stopped ones).
Stop	<code>docker stop <name></code>	Gracefully shuts down the application and container.

Pro Tip: You don't have to type the whole Container Name or ID. Usually, typing the first 3 or 4 characters is enough for Docker to understand which one you mean!



Summary Checklist

1. Write Dockerfile.
2. `docker build .` (To create the image).
3. `docker ps` (To check if things are running).
4. `docker run -p 3000:80 <id>` (To start and connect to your app).

Now that you can run your app, would you like to learn how to give your images "friendly names" (Tags) so you don't have to copy-paste those long IDs anymore?

```
PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup> docker
build .
[+] Building 16.5s (9/9) FINISHED                                            docker:desktop-linux
=> [internal] load build definition from Dockerfile                      0.1s
=> => transferring dockerfile: 147B                                         0.0s
=> [internal] load metadata for docker.io/library/node:latest             0.2s
=> [internal] load .dockerrcignore                                       0.1s
=> => transferring context: 2B                                           0.0s
=> [1/4] FROM docker.io/library/node:latest@sha256:a2f09f3ab9217c692a4e192ea272866ae43b59 0.2s
=> => resolve docker.io/library/node:latest@sha256:a2f09f3ab9217c692a4e192ea272866ae43b59 0.2s
=> [internal] load build context                                         0.1s
=> => transferring context: 304B                                         0.0s
=> CACHED [2/4] WORKDIR /app                                           0.0s
=> [3/4] COPY . /app                                                 0.3s
=> => exporting attestation manifest sha256:06b31a32a72c027946e9eb547f7e9eac0c2      10.0s
1013bad465 0.2s
=> => exporting manifest list sha256:9b8f67c2c1081d84d94cf4df4c4cdbcaad991025d      4.5s
1.8s
PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup> docker run -p 3000:80 160b319d2f1
26e8aeb5aaabfd33adce592971a30f4ae8f9c337644b098a582dd
hi
deep learning
```

```
docker ps          # List running containers
docker logs my-app # View container logs
docker stop my-app # Stop the container
docker rm my-app   # Remove the container
```

It is important to understand that in Docker, you handle **Containers** and **Images** differently.

Because a container is a "running instance," you have to stop it before you can fully remove it.¹

Here is the breakdown of how to shut down and clean up your environment.

1. Stopping and Removing Containers

A container is the "running process" of your app. Even if you close your terminal window, the container might still be running in the background.

Step A: Find the Container

First, you need to know the **Container ID** or the **Name** (the funny name like `wizard_hat` that Docker generates).

Bash

```
docker ps
```

Use `docker ps -a` if you want to see containers that are already stopped.

Step B: Stop the Process

This is like telling the Node.js server to "shut down safely."

Bash

```
docker stop <container_name_or_id>
```

Step C: Remove the Container

Stopping a container doesn't delete it; it just puts it in "Exited" mode. To get rid of it completely:

Bash

```
docker rm <container_name_or_id>
```



2. Removing Images

An image is the "blueprint." You cannot remove an image if a container (even a stopped one) is still using it.

Step A: List your Images

Bash

```
docker images
```

Step B: Remove the Image

Once all containers based on this image are deleted, run:

Bash

```
docker rmi <image_id>
```

(Note: `rmi` stands for "Remove Image".)



3. The "Clear All" Shortcut (Pruning)

If you have been experimenting and have a lot of "trash" (stopped containers and unused images) cluttering your system, you can use the **Prune** command.

- **To remove ALL stopped containers:**
- Bash

docker container prune

-
-
- **To remove ALL unused images, containers, and networks:**
- Bash

docker system prune

-
- *(Use this with caution! It deletes everything that isn't currently running.)*



Summary Table

Goal	Command
Stop a running app	<code>docker stop <id></code>
Delete a container	<code>docker rm <id></code>
Delete an image	<code>docker rmi <image_id></code>
Check what is running	<code>docker ps</code>

Would you like to learn how to run a container in "Detached Mode" so it runs in the background without taking over your terminal screen?

```
PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup> docker images
IMAGE ID          DISK USAGE   CONTENT SIZE   EXTRA
PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup>
p> docker ps
CONTAINER ID     IMAGE          COMMAND          CREATED        STATUS        PORTS
                NAMES
743a43c64a9d   160b319d2f12   "docker-entrypoint.s..."  23 minutes ago  Up 23 minutes  0.0.0.0:3
000->80/tcp, [::]:3000->80/tcp  mystifying_hoover
PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup> docker stop mystifying_hoover
mystifying_hoover
PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup> docker ps
CONTAINER ID     IMAGE          COMMAND          CREATED        STATUS        PORTS
                NAMES
PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup>
```

EXPOSE & A Little Utility Functionality

In the last lecture, we started a container which also **exposed** a port (port 80).

I just want to clarify again, that `EXPOSE 80` in the `Dockerfile` in the end is **optional**. It **documents** that a process in the container will expose this port. But you still need to then actually expose the port with `-p` when running `docker run`. So technically, `-p` is the **only required part** when it comes to listening on a port. Still, it is a **best practice** to also add `EXPOSE` in the `Dockerfile` to document this behavior.

As an additional **quick side-note**: For **all docker commands** where an **ID** can be used, you **don't always have to copy** / write out the **full id**.

You can also **just use the first (few) character(s)** - just enough to have a unique identifier.

So instead of

```
docker run abcdefg
```

you could also run

```
docker run abc
```

or, if there's no other image ID starting with "a", you could even run just:

```
docker run a
```

This applies to ALL Docker commands where IDs are needed.

This is a "lightbulb moment" for most Docker students. You've hit on the **Immutable** (unchangeable) nature of Docker images. Once an image is built, it is essentially a "frozen" snapshot of your project.

Here are the notes on **Image Immutability** and why your code changes didn't show up.



Images are Read-Only (Immutable)

When you run `docker build`, Docker follows the instructions in your `Dockerfile` and creates a static package.

- **The Snapshot:** The `COPY . /app` command takes your code *at that exact second* and bakes it into the image layers.
- **Disconnected:** After the build is finished, there is **no connection** between the files on your hard drive and the files inside the image.



The "Exclamation Mark" Experiment

You found that:

1. **Changing code** on your host machine does nothing to a running container.
 2. **Restarting a container** does nothing because it's still using the "old" image snapshot.
 3. **The Solution:** You must **rebuild** the image to take a *new* snapshot of the updated code.
-



The Rebuild Workflow

When you change your code, the current workflow is:

1. **Modify** your code (e.g., add `!`).
2. **Stop** the old container (`docker stop`).
3. **Rebuild** the image (`docker build .`).
4. **Run** a new container from the **new** Image ID.

Crucial Rule: An image is a "closed template." If you change the ingredients, you have to bake a new cake.

Layer Caching: How Docker Rebuilds Fast

You might have noticed that the second time you ran `docker build ..`, it was much faster than the first time. This is because of **Layer Caching**.

- Docker looks at each instruction in your `Dockerfile`.
- If the instruction and the files it touches haven't changed, Docker reuses the "cached" layer from the previous build.
- As soon as it hits a change (like your updated `server.js`), it rebuilds that layer and **every layer that follows it**.



Looking Ahead: "Bind Mounts"

As you mentioned, rebuilding an image every time you change a single exclamation mark is very slow for development.

Later, you will learn about **Bind Mounts**. This is a way to "link" a folder on your computer directly to a folder inside the container. This allows the container to see code changes **instantly** without needing a rebuild. But for now, understanding that images are static is the most important foundation.

Would you like to explore why the `docker build` command was faster the second time by looking at how Docker uses its internal cache?

This is one of the most important lessons in Docker "mastery." You've moved from just making things work to making them **efficient**. Understanding that a `Dockerfile` is a stack of layers—and that the **order** of those layers matters—is what separates beginners from pros.

Here are your notes on **Layer Caching** and **Image Optimization**.



The Layer-Based Architecture

Every single instruction in your `Dockerfile` (FROM, RUN, COPY, etc.) creates a **layer**. An image is just a stack of these read-only layers.

1. How Caching Works

When you run `docker build`, Docker looks at your instructions:

- If an instruction hasn't changed (and the files it touches haven't changed), Docker says "**Using cache**" and skips the work.

- **The Domino Effect:** As soon as one layer changes (like a code edit), that layer **and every layer below it** must be rebuilt. Docker cannot "skip" a middle layer if the one above it is new.

2. The "Running" Layer

When you finally run the container, Docker adds a thin **Read/Write Layer** (the "Container Layer") on top of the image's "Read-Only" layers. This is where the actual application lives and breathes while it's running.

⚡ Optimization: The "Package First" Pattern

In your original Dockerfile, `npm install` ran every time you changed a single letter in `server.js`. This is because `COPY . /app` was placed *before* `RUN npm install`.

Since `.` includes your source code, any code change invalidated the `COPY` layer, which forced the `RUN npm install` layer to trigger again.

The Optimized Dockerfile

By splitting the copy process, we "protect" the expensive `npm install` step from code changes.

Dockerfile

```
FROM node
```

```
WORKDIR /app
```

```
# Step A: Copy ONLY the dependency list
COPY package.json /app
```

```
# Step B: Install dependencies (This is now cached!)
RUN npm install
```

```
# Step C: Copy the rest of the source code
COPY . /app
```

```
EXPOSE 80
```

```
CMD ["node", "server.js"]
```

Why this is better:

1. **If you change `server.js`:** Docker sees `package.json` is the same, so it uses the cache for `npm install`. It only re-executes the final `COPY . /app` (which is near-instant).

2. **If you change package.json:** Docker realizes you have new dependencies, so it re-runs `npm install` to update the image.
-



Core Takeaway

Order matters. Always put your **least frequently changed** instructions at the top and your **most frequently changed** instructions (like your source code) at the bottom.

Would you like to see how to "Tag" your images now? Instead of using those long IDs, we can name this optimized version something like goals-app:v2-optimized.

docker management commands

This section highlights a major shift in how you use Docker day-to-day: moving from **creating** containers to **managing** them.

The biggest takeaway here is that `docker run` and `docker start` are **not** the same thing.



Managing the Lifecycle: Run vs. Start

Understanding the difference between these two commands will save you from accidentally creating hundreds of "ghost" containers on your hard drive.

1. `docker run` (The Creator)

Think of this as "Create + Start."

- It takes an **Image**.
- It creates a brand-new **Container** (a fresh box).
- It **Starts** that container.
- **Problem:** If you use `docker run` every time you want to start your app, you will end up with a new container every single time, which eats up disk space.

2. `docker start` (The Restarter)

Think of this as "Turn on the power."

- It takes an **existing Container ID or Name**.
- It restarts a container that was previously stopped.
- It preserves the container's state (any files changed inside stay there).
- It runs in **Background Mode** (Detached) by default.



Essential Management Commands

Command	Usage	Description
<code>docker --help</code>	<code>docker --help</code>	The "Manual." Shows every available command.
<code>docker ps</code>	<code>docker ps</code>	Lists only currently running containers.

<code>docker ps -a</code>	<code>docker ps -a</code>	Lists all containers (Running + Stopped).
<code>docker start</code>	<code>docker start <name></code>	Powers on a stopped container (runs in background).

Understanding "Detached" Mode

You noticed that when you used `docker start`, it didn't "block" your terminal like `docker run` did. This is a core Docker behavior:

- **Attached Mode (Default for run):** You see the logs/output from the container directly in your terminal. You can't type new commands until you stop the container.
- **Detached Mode (Default for start):** The container runs "in the background." You get your terminal back immediately to keep working.

Verification: Since the terminal is free, you can't see the Node.js logs. You have to verify the app is alive by running `docker ps` or refreshing your browser at `localhost:3000`.

Quick Tip: Help for Subcommands

As you noted, you can get specific help for any subcommand. If you are confused about how to filter your process list, just run:

```
docker ps --help
```

Now that your container is running in the background, would you like to learn how to "attach" back to it so you can see the logs, or perhaps how to run a container in detached mode from the very beginning?

```

PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup> docker
ps -a
● CONTAINER ID   IMAGE          COMMAND           CREATED          STATUS          P
● 743a43c64a9d   160b319d2f12   "docker-entrypoint.s..."   2 hours ago    Exited (137) 2 hours ago
●      mystifying_hoover
PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup> docker
start mystifying_hoover
mystifying_hoover
● PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup> docker
● ps
CONTAINER ID   IMAGE          COMMAND           CREATED          STATUS          PORTS
NAMES
743a43c64a9d   160b319d2f12   "docker-entrypoint.s..."   2 hours ago    Up 10 seconds   0.0.0.0:3000
->80/tcp, [::]:3000->80/tcp   mystifying_hoover
○ PS C:\Users\Public\docker-learning\first-demo-starting-setup-1\nodejs-app-starting-setup>

```

This section tackles a very common point of confusion: "**Where did my console logs go?**" When you run an app, you often need to see the output for debugging. Docker gives you full control over whether you want to "watch" the app or let it run quietly in the background.

Attached vs. Detached Mode

The "Mode" refers to whether your terminal is currently linked to the container's output (logs) and input.

1. Attached Mode (Default for `docker run`)

Your terminal is "locked" to the container.

- **Pros:** You see every `console.log()` or error message immediately.
- **Cons:** You can't use that terminal for anything else until the container stops.

2. Detached Mode (Default for `docker start`)

The container runs in the background.

- **Pros:** Your terminal is free. You can run more commands or start other containers.
- **Cons:** You don't see what's happening inside the "box" by default.

The "Mode" Command Cheat Sheet

You are not stuck with the defaults! You can swap modes using these flags and commands:

To do this...	Use this command
Run in the background	<code>docker run -d <image-id></code>
Restart and watch logs	<code>docker start -a <container-name></code>
Re-join a running app	<code>docker attach <container-name></code>
See past logs	<code>docker logs <container-name></code>
Watch logs live	<code>docker logs -f <container-name></code>

Deep Dive: `docker logs` vs `docker attach`

If a container is already running in detached mode, you have two ways to see what it's doing:

A. The Snapshot: `docker logs`

Running `docker logs <name>` is like taking a screenshot of the terminal history. It shows you everything that has happened *up until now* and then finishes.

B. The Live Stream: `docker logs -f`

The `-f` stands for **Follow**. This keeps the terminal open and waits for new logs to appear. This is often better than `attach` because it is easier to exit without accidentally stopping the container.



Key Takeaway

- `docker run -d` is the standard for production or stable environments where you don't need to watch every move.
- `docker logs -f` is your best friend during development when you are testing if your code (like the "Goal" setter) is actually working.

Would you like to learn how to enter "Interactive Mode" inside a custom image, similar to how we used the `-it` flag with the official Node image earlier?

Attaching to an already-running Container

By default, if you run a Container without `-d`, you run in "attached mode".

If you started a container in detached mode (i.e. with `-d`), you can still attach to it afterwards without restarting the Container with the following command:

```
docker attach CONTAINER
```

attaches you to a running Container with an ID or name of `CONTAINER`.

This section is a great shift in perspective. It proves that Docker isn't just for background web servers—it's for any application, including **interactive command-line tools**.

You've highlighted a critical realization: some apps *require* a two-way conversation between the user and the container.



Dockerizing a Python Utility

Unlike the Node.js web server that stays "open" waiting for web requests, this Python script is a **Utility App**. It follows a specific lifecycle:

1. Start → 2. Ask for input → 3. Process → 4. Output → 5. Finish/Exit.

The Python Dockerfile

The structure remains familiar, showing how consistent Docker is across languages:

```
Dockerfile
```

```
FROM python
```

```
WORKDIR /app
```

```
COPY . /app
```

```
# No EXPOSE needed because there is no web traffic!
```

```
CMD ["python", "rng.py"]
```

The Interactive "Missing Link": `-it`

When you first ran the Python container, it crashed or hung. This is because, by default, a Docker container doesn't "listen" to your keyboard; it only shows you what the app "says" (STDOUT).

To fix this, you used the "Golden Duo" of flags:

Flag	Name	Function
<code>-i</code>	Interactive	Keeps the input stream (STDIN) open. The container is now "listening" for your typing.
<code>-t</code>	TTY	Allocates a "pseudo-terminal." This makes the container look and behave like a real terminal window.

Command: `docker run -it <image-id>`

Restarting with Interaction

As you discovered, `docker start` defaults to detached mode. If you want to restart an interactive utility, you have to explicitly tell Docker to re-open the communication lines.

- `docker start -i <name>`: Restarts the container and lets you type into it (Interactive).
- `docker start -a -i <name>`: Restarts, lets you type, **and** shows you the output (Attached + Interactive).

Key Takeaways

1. **Docker is Universal:** It doesn't care if it's running a complex Node.js API or a 10-line Python script.
 2. **Standard Input (STDIN):** For any app that asks "What is your name?" or "Enter a number," you **must** use the `-i` flag.
 3. **The Lifecycle:** Utility containers stop automatically when the script finishes. This is different from web servers, which run until you manually stop them.
-

Would you like to learn about "Cleanup on Exit"? Since these utility containers stop as soon as they are done, I can show you a flag that deletes them automatically so they don't clutter your docker ps -a list.

delete the container, images

This lesson focuses on **Cleanup**. Docker is amazing for creating environments, but it can quickly clutter your hard drive with "zombie" containers and large images. Knowing how to delete what you don't need is essential for a healthy development workflow.



Managing Containers: rm

When a container stops, it doesn't disappear; it stays on your disk in an "Exited" state.

1. Manual Removal

- **Command:** docker rm <container_name_or_id>
- **Multiple:** You can remove several at once by listing them: docker rm con1 con2 con3.
- **The Rule:** You **cannot** remove a running container. You must docker stop it first.

2. Automatic Cleanup: The --rm Flag

As you hinted, there is a much more elegant way. If you know you only need a container temporarily (like the Python RNG utility), add the `--rm` flag when you start it:

- **Command:** docker run -it --rm <image_id>
 - **Benefit:** As soon as the container stops (either because the app finished or you stopped it), Docker **automatically deletes** it. No more messy ps -a lists!
-



Managing Images: rmi

Images are the "blueprints." Because they contain entire operating systems (like Linux) and environments (like Node), they can be quite large.

1. Manual Removal

- **Command:** docker rmi <image_id>
- **The Dependency Rule:** You cannot remove an image if it is being used by a container.
 - This includes **stopped** containers!
 - **Workflow:** Delete the Container (rm) → Delete the Image (rmi).

2. Deep Dive: Image Size

You noticed that your custom image is almost the same size as the official node image.

- **Layers:** This is because your image **is** the Node image, plus a tiny layer of your code.
 - **Shared Space:** Docker is smart. If you have 10 images based on node, Docker only stores the actual node files **once** on your hard drive to save space.
-

The "Power Wash": `prune`

If your system feels slow or your disk is full, you can use the `prune` commands to clear out the "junk."

Command	What it deletes
<code>docker container prune</code>	Deletes all stopped containers.
<code>docker image prune</code>	Deletes all "dangling" images (images without tags).
<code>docker image prune -a</code>	Deletes all images not used by at least one container.
<code>docker system prune</code>	The ultimate cleanup. Deletes all stopped containers, unused networks, and dangling images.

Summary Table: RM vs RMI

Target	Command	Requirement
Container	<code>docker rm</code>	Must be stopped first.
Image	<code>docker rmi</code>	All associated containers must be deleted first.

Now that our system is clean, would you like to learn how to name (Tag) your images so that instead of typing `docker rmi 7a2b9...`, you can type `docker rmi my-node-app:v1`?

This lesson introduces one of the most practical "quality of life" flags in the Docker CLI: the `--rm` flag. As your projects grow, your system can quickly fill up with "dead" containers that have served their purpose. This flag automates the cleanup.



Automatic Cleanup with `--rm`

In the standard Docker workflow, when a container stops, it switches to an "Exited" state. It still exists on your disk, taking up resources, until you manually run `docker rm`.

How it Works

Adding `--rm` to your `docker run` command tells Docker: "*As soon as the process inside this container stops, delete the container immediately.*"

The Command:

Bash

```
docker run -p 3000:80 -d --rm <image-id>
```



When to use `--rm`

While you can use it for almost anything, it is especially useful in two scenarios:

1. Web Servers (like your Node.js app)

Since a web server container is often only stopped when you've changed your code, you'll likely need to rebuild a new image and start a fresh container anyway. There is rarely a reason to keep the old, stopped "version" of that container lying around.

2. Utility Tools (like your Python RNG script)

For apps that perform a quick task and then finish, `--rm` is perfect. It ensures that your `docker ps -a` list stays clean and only shows what you actually care about.



The Full "Pro" Command

By combining everything you've learned so far, your standard command for running a web application likely looks like this:

Flag	Purpose

<code>-p 3000:80</code>	Publish: Connects your computer's port 3000 to the container's port 80.
<code>-d</code>	Detached: Runs the app in the background so you can keep using your terminal.
<code>--rm</code>	Remove: Deletes the container automatically once you stop it.



Core Takeaway

The `--rm` flag helps you follow the "**Cattle, not Pets**" philosophy of DevOps. Containers are temporary and replaceable. If you need to change something, you kill the old one and start a new one—`--rm` just handles the "trash" for you.

Now that you have mastered running and cleaning up containers, would you like to learn about "Inspecting" an image? I can show you how to see the "metadata" of an image (like its history and configuration) without actually running it.

inspect images

This section dives into the "DNA" of your Docker images. While the Dockerfile is the source code for your image, the **Inspect** command allows you to see the final, compiled metadata that Docker uses to run your containers.



Peeking Inside: docker image inspect

The inspect command provides a deep dive into the configuration of an image. It returns a JSON object containing everything Docker knows about that specific image.

Command: docker image inspect <image-id>



What can you find in the metadata?

When you run this command, you'll see a long list of properties. Here are the key highlights:

- **ID & Created:** The full cryptographic hash of the image and the exact timestamp it was built.
 - **ContainerConfig:** This is the most important part. It shows the **Environment Variables**, the **Command (CMD)**, and the **Exposed Ports** that were "baked" into the image.
 - **OS/Architecture:** Confirms which operating system the image is based on (usually a flavor of Linux) and the CPU architecture (like `amd64` or `arm64`).
 - **Layers:** A list of the filesystem changes that make up the image.
-



The Relationship: Image vs. Container Layers

It is important to visualize how your computer handles memory and storage with Docker.

1. **Shared Foundation:** The image layers are **Read-Only**.
 2. **Efficiency:** If you run 5 containers from the same Node.js image, Docker **does not** copy the code 5 times. All 5 containers point to the exact same read-only image layers on your disk.
 3. **The Writeable Layer:** When you start a container, Docker adds a thin "**Container Layer**" on top. If your app creates a log file or a user uploads a photo, that data is stored *only* in this thin, top layer.
-



Understanding the Layer Count

You noticed that your `inspect` output showed more layers than you had instructions in your Dockerfile. This is because layers are **cumulative**:

- **Your Layers:** `WORKDIR`, `COPY`, `RUN`, etc.
- **Base Image Layers:** The official `node` image you used with `FROM` already has its own layers (Linux OS + Node.js installation).

- **The Chain:** When you build your image, you are simply adding your layers on top of the existing stack from the base image.
-



Why use inspect?

You won't use this every day, but it is a lifesaver in these situations:

- **Mystery Images:** You pulled an image from Docker Hub and want to know which port it expects you to use.
 - **Debugging:** You want to verify if your Environment Variables or `WORKDIR` were set correctly during the build.
 - **Storage Management:** You want to see exactly how many layers an image has and how they are structured.
-

Now that you can look "under the hood" of an image, would you like to learn how to copy files specifically between your host machine and a running container without using the `COPY` instruction?

`docker cp`

This lecture introduces a handy "utility" command that breaks the standard rule of container isolation. While we usually think of containers as closed boxes, the `cp` (copy) command gives you a direct way to move files in and out of a running container without needing to rebuild or restart anything.

The docker cp Command

The `cp` command is the bridge between your **Host Machine** (your computer) and the **Container File System**.

1. Copying from Host → Container

This is useful if you want to drop a configuration file or a temporary script into a container that is already running.

Syntax:

```
docker cp <local-path> <container-name>:<path-inside-container>
```

- **Example:** `docker cp ./dummy/. my_container:/app/test`
- **Result:** All files inside your local `dummy` folder are moved to the `/app/test` folder inside the container.

2. Copying from Container → Host

This is arguably the more powerful use case. Since containers are "black boxes," extracting data is often necessary for debugging.

Syntax:

```
docker cp <container-name>:<path-inside-container> <local-path>
```

- **Example:** `docker cp my_container:/app/logs.txt ./local_logs`
- **Result:** You "pull" a file out of the container so you can open it in your favorite text editor on your computer.

Use Cases: Why use cp?

Use Case	Why it's helpful

Log Extraction	If your app crashes or generates internal logs, you can pull them out to analyze them locally.
Quick Fixes	Changing a single line in a <code>.conf</code> file to see if it fixes a bug without doing a full 2-minute build/run cycle.
Data Recovery	If a container is about to be deleted but contains a file you need to save, you can rescue it with <code>cp</code> .

⚠️ Important Limitations

While `cp` is powerful, the video highlights why you **shouldn't** use it for everything:

1. **Error-Prone for Code:** Updating your source code using `cp` is dangerous. It's easy to miss a file, leading to a "Frankenstein" app where some files are new and some are old.
2. **Execution Locks:** You often cannot replace a file that is currently being executed (like `your-server.js` while the Node server is running).
3. **Temporary:** Any files you copy **into** a container will be lost forever if that container is deleted. Remember, containers are temporary; images are permanent!

Preview: If you are tired of rebuilding images to see code changes, we will eventually learn about **Volumes** and **Bind Mounts**, which act like a "live sync" between your computer and the container.

🧠 Summary Checklist

- **Move to Container:** `docker cp [source] [container]:[dest]`
- **Move to Host:** `docker cp [container]:[source] [dest]`
- **State:** The container can be running or stopped.
- **Creation:** Docker will automatically create the destination folder if it doesn't exist yet.

Now that you can move files around, would you like to learn how to "Rename" your containers or images? This makes commands like `docker cp` much easier since you won't have to use those random generated names like `boring_wozniak`.

naming to Images and containers

This final piece of the puzzle makes your Docker workflow feel professional and organized. Instead of memorizing random strings like `7a2b9...` or `wizard_hat`, you can now use human-readable names for both your **Containers** and your **Images**.



Naming Containers: The `--name` Flag

By default, Docker assigns a random name (e.g., `focused_curi`). To take control, you use the `--name` flag during the `run` command.

- **Command:** `docker run -d --rm -p 3000:80 --name my-app goals:latest`
 - **Advantage:** You can now stop or delete the container without ever running `docker ps`. You just know the name: `docker stop my-app`.
-



Tagging Images: The "Name:Tag" Structure

Image names are actually called **Tags**, and they follow a two-part format: `repository:tag`.

1. The Repository (The "Group" Name)

This is the general name of your application, like `node`, `python`, or `goals-app`.

2. The Tag (The "Version" or "Flavor")

This is the specific version or configuration. If you don't provide a tag, Docker defaults to `:latest`.

- **Versions:** `node:14`, `node:12`, `goals:v1.0`
 - **Flavors:** `node:slim` (a smaller, lightweight version of the image).
-



How to Tag Your Images

You usually tag an image at the moment you build it using the `-t` (or `--tag`) flag.

Command: `docker build -t goals:latest .`

Why use Tags?

1. **Version Control:** You can keep an image for `v1` and a new image for `v2` on your machine at the same time.
 2. **Environment Matching:** As you saw, using `FROM node:12` ensures your container uses the exact same version of Node.js that your production server uses, preventing the "it works on my machine" bug.
 3. **Readability:** `docker run goals:v1` is much clearer than `docker run 4c12...`
-



Summary of Naming Commands

Feature	Concept	Command
Container Name	Friendly identifier for a process.	<code>docker run --name my-container <image></code>
Image Repository	The "App Name."	<code>docker build -t app-name .</code>
Image Tag	The "Version Name."	<code>docker build -t app-name:v1 .</code>
Official Tags	Specialized base images.	<code>FROM node:14-slim</code>



Final Review Checklist

You have now covered the "Core Essentials" of Docker:

- [x] **Images vs. Containers:** Blueprints vs. Running Instances.
- [x] **Dockerfile:** The recipe for custom images.
- [x] **Layer Caching:** Optimizing build speed by ordering instructions.
- [x] **Port Mapping:** Connecting your PC (-p) to the container.
- [x] **Interactive Mode:** Using -it for utility apps.
- [x] **Cleanup:** Using --rm, prune, and rmi to keep your system lean.
- [x] **Tags & Names:** Organizing your work like a pro.

You have a solid foundation! Would you like me to create a "Cheat Sheet" summarizing all the commands we've covered in these notes for your future reference?

share images-docker hub

Sharing images is the "killer feature" of Docker. You've perfectly summarized the two philosophies: sharing the **recipe** (Dockerfile) versus sharing the **prepared meal** (the built Image).

Here are the organized notes on why and how we move from local development to a shared ecosystem.



Why Share? The "Clean Machine" Advantage

As you pointed out, Docker solves the "**Matrix of Hell**"—where different projects require different versions of the same tool.

- **Version Coexistence:** You can run a Python 3.7 app and a Python 3.9 app simultaneously. Neither "knows" the other exists, and neither messes with your global Windows/Mac/Linux settings.
- **Zero Installation:** A new team member doesn't need to spend a day installing Node, Python, or Databases. They just need Docker.



The Two Paths of Sharing

1. Sharing the Dockerfile (The "Recipe")

This is what you find in the course attachments. You provide the **Dockerfile + Source Code**.

- **Pros:** The recipient can see exactly how the image is built and modify the code.
- **Cons:** The recipient **must** run the `docker build` step. This takes time and requires them to have all the source files locally.
- **Best for:** Development and collaboration where code changes are expected.

2. Sharing the Built Image (The "Finished Product")

This is how official images like `node` or `python` are shared. You provide the **compiled, multi-layered image**.

- **Pros: No build step.** The recipient runs `docker run` and the app starts instantly. They don't even need the source code files; everything is already "baked" into the layers.
- **Cons:** The recipient can't easily see or change the source code (it's a "black box").
- **Best for:** Deployment, production, and using standard tools (like a Database or Web Server).



Understanding the Base Image (The `node` Example)

You made an excellent point about the `node` Dockerfile on GitHub. Even though we use `FROM node`, we aren't downloading that Dockerfile from GitHub and building it.

1. **The Node Team** wrote a complex Dockerfile (based on a tiny Linux OS called **Alpine**).
 2. **The Node Team** built that image and uploaded the "finished product" to **Docker Hub**.
 3. **You** simply "pull" that finished product. Your machine downloads the layers, and you start your work from Step 10 instead of Step 0.
-



What's Next?

To share a "Finished Image" like the Node team does, we need a place to put it. This brings us to **Docker Registries** (like Docker Hub).

Would you like to learn how to "Push" your custom goals-app image to Docker Hub so that you (or anyone else) can download it on a completely different computer?

push-rename-login

This is the final step in turning your local project into a globally available tool. By pushing to **Docker Hub**, you are moving your "finished meal" from your local kitchen to a public restaurant where anyone can "order" (pull) it.

Here are the notes on the **Docker Hub Workflow**.



The Docker Hub Workflow

To share your image, you follow a specific 4-step process: **Create, Tag, Login, and Push**.

1. Create the Repository (on Hub)

Before you can push, you must have a "parking space" ready on Docker Hub.

- Log in to hub.docker.com.
- Create a new repository (e.g., `node-hello-world`).
- Note your full repository path: `your-username/node-hello-world`.

2. The "Naming" Rule: Re-Tagging

Docker Hub is strict. You cannot push an image named simply `node-demo`. The local image name **must** exactly match your Docker Hub repository path.

If you already built the image with the "wrong" name, you don't need to rebuild it. Use the `tag` command to create a pointer (clone):

- **Command:** `docker tag node-demo:latest your-username/node-hello-world:latest`
- **Result:** You now have two names pointing to the exact same image ID.

3. Authentication: `docker login`

To prove you have permission to push to your account, you must log in via your terminal.

- **Command:** `docker login`
- **Note:** When typing your password, the terminal will stay blank (no asterisks). This is a standard security feature.

4. The Push: `docker push`

Now you can upload your image layers to the cloud.

- **Command:** `docker push your-username/node-hello-world:latest`



Efficiency: Layered Uploads

You noticed that the upload was surprisingly fast. This is because of **Layer Reuse**:

- Docker Hub already has the official node layers.
 - When you push your image, Docker Hub says: "Wait, I already have the 900MB of Node.js layers. Just send me your 2MB of custom code!"
 - This saves bandwidth for you and storage space for Docker.
-



How Others Use Your Image

Once your image is on Docker Hub, a teammate (or you on a different computer) can run it instantly:

1. They **don't** need your source code.
2. They **don't** need your Dockerfile.
3. They simply run:
`docker run your-username/node-hello-world`



Summary of Commands

Action	Command
Authenticate	<code>docker login</code>
Rename/Clone	<code>docker tag <old-name> <new-name></code>
Upload	<code>docker push <username>/<repo>:<tag></code>
Download	<code>docker pull <username>/<repo>:<tag></code>

Congratulations! You have successfully completed the entire Docker lifecycle.

Would you like to wrap up this module with a **Master Cheat Sheet** of all the commands we've used, or are you ready to dive into the next big topic: **Managing Data with Volumes and Bind Mounts?**

pull

This perfectly illustrates the final loop of the Docker ecosystem: **The Pull and Run**. You've touched on the most important behavior of the Docker CLI—how it decides whether to download something or use what it already has.



The "Pull" Behavior

When you share a public repository on Docker Hub, you are essentially creating a global "app store" for your code.

1. Public Accessibility

As you demonstrated, **Authentication is one-way** for public repos:

- **Pushing:** Requires `docker login` (only the owner can change the image).
- **Pulling:** Anyone can do it without logging in. This is why official images like `node` and `python` are so easy to use.

2. Explicit vs. Implicit Pulling

There are two ways to get a remote image onto your machine:

- `docker pull <name>`: This is a manual "Download" button. It checks the registry and brings the image layers to your local storage.
 - `docker run <name>`: If Docker doesn't find the image locally, it acts as an **Implicit Pull**. It will automatically search Docker Hub, download it, and then start the container.
-



The "Local First" Trap (Crucial!)

This is the most common mistake developers make when working with teams or deployment.

Docker is "lazy" by design. If you run `docker run my-app`, Docker follows this logic:

1. Does `my-app` exist on this computer?
2. **Yes?** Run it immediately. (It does **not** check if there is a newer version on Docker Hub).
3. **No?** Go to Docker Hub, download it, and then run it.

The Solution: Manual Updates

If your teammate pushes a fix to `academind/node-hello-world:latest`, and you already have an older version of that image on your machine, simply typing `docker run` will run your **old** version.

To get the update, you must force the sync:

1. `docker pull academind/node-hello-world:latest` (This replaces your local version with the one from the cloud).
2. `docker run ...` (Now uses the updated image).



Summary of Sharing & Pulling

Command	Action	Local Check?
<code>docker pull</code>	Fetch latest version from Hub.	Always checks Hub.
<code>docker run</code>	Create & Start container.	Uses local if it exists; pulls if it doesn't.
<code>image prune -a</code>	Delete unused images.	Useful for "cleaning the slate" before a pull.

🎓 Module Wrap-Up: You are now a Docker Navigator!

You have mastered the core cycle:

1. **Create** (Dockerfile)
2. **Build** (Image)
3. **Run** (Container)
4. **Optimize** (Layers)
5. **Manage** (Cleanup/Names/Tags)
6. **Share** (Push/Pull)

What's next? Now that you know how to move images around, the next big challenge is **Data**. Currently, if you delete a container, your data (like your "Goals") disappears.

Would you like to start the next module on "Data & Volumes" to learn how to make your data survive even when containers are deleted?

Tab 12

docker-springboot

springboot-docker

Dockerfile :

```
-----  
FROM eclipse-temurin:17  
  
LABEL maintainer="bhuvandosapati517@gmail.com"  
  
WORKDIR /app  
  
COPY target/docker-learning-demo-0.0.1-SNAPSHOT.jar  
/app/springboot-docker-demo.jar  
  
ENTRYPOINT ["java","-jar","springboot-docker-demo.jar"]
```

commands

Id CommandLine

```
-- -----
1 docker build -t springboot-demo-docker .
4 docker images
5 docker build -t springboot-demo-docker:0.1.RELEASE .
7 docker run -p 8080:8080 springboot-demo-docker:0.1.RELEASE
10 docker run -p 8080:8080 -d springboot-demo-docker:0.1.RELEASE
11 docker ps -a

13 docker stop condescending_greider
14 docker ps

16 docker run -p 8081:8080 -d springboot-demo-docker:0.1.RELEASE
18 clear
21 docker rm condescending_greider
22 docker rm silly_hertz
23 docker stop silly_hertz
25 docker images
26 docker login
27 docker tag springboot-demo-docker:0.1.RELEASE bhuvan0517/springboot-demo-docker

29 docker tag springboot-demo-docker bhuvan0517/springboot-demo-docker:0.1.RELEASE
30 docker images
31 docker rmi bhuvan0517/springboot-demo-docker:latest
32 docker images
34 docker push bhuvan0517/springboot-demo-docker:0.1.RELEASE

36 docker pull bhuvan0517/springboot-demo-docker:0.1.RELEASE

38 docker pull mysql:latest
39 docker images
40 docker run -p 3306:3306 --name localhost -e MYSQL_ROOT_PASSWORD=Mysql@123 -e
MYSQL_DATABASE=employee_db -e MYSQL_USER=bhuvan -e MYSQL_PASSWORD=bhuvi
...
41 docker images
42 docker ps
43 docker logs -f mysql:latest
44 docker logs -f localhost
45 clear
46 docker exec -it localhost bash
```

